Cloud DevOps Runbook: Incident Management

Document ID: OPS-RUNBOOK-2026-v1 | Google Cloud Run, Vertex AI & Firestore | SRE Team

1. Production Architecture & Sizing

cortex.ai workloads run containerized on Google Cloud Run with serverless scale-to-zero economics. Production guardrails restrict maximum concurrent instances to 2, ensuring budget protection within the $300 GCP credit allocation.

2. Change Windows & Release Freezes

- **Standard Release Windows:** Tuesdays & Thursdays (10:00 AM – 2:00 PM UTC).
- **Weekend Freeze:** Production deployments are frozen from Friday 12:00 PM UTC through Monday 8:00 AM UTC.
- **Quarter-End Freeze:** No releases during the final 5 business days of each fiscal quarter.

3. High-Risk Mutations & Human-in-the-Loop Safeguards

Infrastructure restarts, database migrations, and IAM modifications cannot be performed autonomously. cortex.ai automatically generates an amber **Pending Confirmation Action Card** that holds the execution until human DevOps sign-off is confirmed.

4. Emergency Cloud Run Traffic Rollback

If 5xx errors exceed 1.0%, revert traffic immediately using:
gcloud run services update-traffic cortex-hub --to-revisions=PREVIOUS=100